

The magic of WITH: Common Table Expressions

Jim Nasby, OpenSCG

<https://github.com/decibel/presentations/blob/master/CTEs.pdf>

- What NOT to do with CTEs
- CTEs for chaining DML
- Recursive CTEs

What not to do

- CTEs ***materialize*** everything
- WHERE clauses will not be “pushed down” to a CTE
- This means using CTEs for code clarity can seriously hurt performance

Bad

```
WITH
  tables AS (
    SELECT nspname AS table_schema, relname AS table_name
      , c.oid AS relid
    FROM pg_namespace n
      JOIN pg_class c ON n.oid = c.relnamespace
    WHERE relkind = 'r'
  )
, columns AS (
  SELECT table_schema, table_name, attname AS column_name
    , typename AS column_type
  FROM pg_attribute a
    JOIN pg_type type ON type.oid = atttypeid
    JOIN tables tab ON tab.relid = attrelid
  )
SELECT * FROM columns WHERE table_name = 'pg_class'
;
```

Bad

QUERY PLAN

```
-----  
CTE Scan on columns  
  Filter: (table_name = 'pg_class'::name)  
CTE tables  
  -> Hash Join  
        Hash Cond: (c.relnamespace = n.oid)  
  -> Seq Scan on pg_class c  
        Filter: (relkind = 'r'::"char")  
        -> Hash  
              -> Seq Scan on pg_namespace n  
CTE columns  
  -> Hash Join  
        Hash Cond: (a.atttypid = type.oid)  
        -> Hash Join  
              Hash Cond: (a.attrelid = tab.relid)  
            -> Seq Scan on pg_attribute a  
                  -> Hash  
                        -> CTE Scan on tables tab  
                  -> Hash  
                        -> Seq Scan on pg_type type  
(19 rows)  
Execution time: 5.693 ms
```


Good

```
SELECT table_schema, table_name
      , column_name, column_type
FROM (
    SELECT nspname AS table_schema, relname AS table_name
          , c.oid AS relic
    FROM pg_namespace n
         JOIN pg_class c ON n.oid = c.relnamespace
    WHERE relkind = 'r'
) AS tables
JOIN (
    SELECT attname AS column_name, typename AS column_type
          , attrelid
    FROM pg_attribute a
         JOIN pg_type type ON type.oid = atttypid
) AS columns
ON tables.relid = columns.attrelid
WHERE table_name = 'pg_class'
;
```

Good

QUERY PLAN

Nested Loop

-> Nested Loop

-> Nested Loop

Join Filter: (n.oid = c.relnamespace)

-> **Index Scan using pg_class_relname_nsp_index on pg_class c**

Index Cond: (relname = 'pg_class'::name)

Filter: (relkind = 'r'::"char")

-> **Seq Scan on pg_namespace n**

-> **Index Scan using pg_attribute_relid_index on pg_attribute a**

Index Cond: (attrelid = c.oid)

-> **Index Scan using pg_type_oid_index on pg_type type**

Index Cond: (oid = a.atttypid)

(12 rows)

Execution time: 0.462 ms

Materialization

- CTEs store their results in a “Tuple Store”
- Think temporary table, with no possibility of indexes, but a bit less overhead.

CTEs for DML

CTEs allow use of the RETURNING clause in a DML statement to do multiple things with that data

- INSERT into invoice and invoice_item tables in one statement
- DELETE from one table and INSERT data into another
- Chained UPDATE
- SELECT current data from a table and INSERT it into another table before UPDATEing the same data

Chained INSERT

```
CREATE TABLE invoice(  
    invoice_id serial PRIMARY KEY  
);  
CREATE TABLE invoice_item(  
    invoice_id int REFERENCES invoice  
    , line_number int  
    , item text  
);
```

Chained INSERT

```
WITH invoice AS (  
  INSERT INTO invoice VALUES (default)  
  RETURNING invoice_id  
)  
INSERT INTO invoice_item(invoice_id, ...)  
SELECT invoice_id, line_number, item  
FROM invoice  
  -- Create two lines  
  , ( VALUES  
      (1, 'pizza')  
      , (2, 'soda')  
    ) line(line_number, item)  
RETURNING *  
;
```

Chained INSERT

invoice_id	line_number	item
5	1	pizza
5	2	soda

(2 rows)

Chained INSERT

```
WITH invoice AS(  
  INSERT INTO invoice VALUES (default),(default)  
  RETURNING *  
)  
SELECT *  
FROM invoice  
LIMIT 1  
;
```

invoice_id
6

(1 row)

```
SELECT max(invoice_id) FROM invoice;
```

Chained INSERT

```
SELECT max(invoice_id) FROM invoice;
```

max
7

(1 row)

CTEs always fully execute any DML operations.

Recursive CTEs

Critical to think in set theory

- Start with a single set of rows
- UNION them with a SELECT that's self referencing
- Don't forget the WHERE clause if necessary

Recursive CTEs

What will this output?

```
WITH RECURSIVE s AS (  
  SELECT 1 AS seq  
  UNION ALL  
  SELECT seq + 1  
    FROM s  
   WHERE seq <= 5  
) SELECT * FROM s;
```


Recursive CTEs

```
WITH RECURSIVE s AS (  
  SELECT 1 AS seq  
  UNION ALL  
  SELECT seq + 1  
    FROM s  
  WHERE seq <= 5 -- Applied to input, not output  
) SELECT * FROM s;
```

seq
1
2
3
4
5
6

(6 rows)

Recursive CTEs

```
WITH RECURSIVE s AS (  
    SELECT 1 AS seq  
    UNION ALL  
    SELECT seq + 2  
    FROM s  
    WHERE seq <= 5  
) SELECT * FROM s;
```

Recursive CTEs

```
WITH RECURSIVE s AS (  
  SELECT 1 AS seq  
  UNION ALL  
  SELECT seq + 2  
  FROM s  
  WHERE seq <= 5  
) SELECT * FROM s;
```

seq
1
3
5
7

(4 rows)

Recursive CTEs

```
WITH RECURSIVE s AS (  
    SELECT 1 AS seq  
UNION ALL  
    SELECT seq + 1  
        FROM s  
        WHERE seq <= 5  
) SELECT * FROM s;
```

```
WITH RECURSIVE s AS (  
    SELECT 1 AS seq  
UNION ALL  
    SELECT seq + 1 FROM s  
)  
SELECT *  
    FROM s  
    WHERE seq <= 5  
;
```


Recursive CTEs

QUERY PLAN

CTE Scan on s (cost=2.95..3.57 rows=31 width=4)

CTE s

-> Recursive Union (cost=0.00..2.95 rows=31 width=4)

-> Result (cost=0.00..0.01 rows=1 width=4)

-> WorkTable Scan on s s_1 (cost=0.00..0.23 rows=3 width=4)

Filter: (seq <= 5)

(6 rows)

QUERY PLAN

CTE Scan on s (cost=4.28..6.55 rows=34 width=4)

Filter: (seq <= 5)

CTE s

-> Recursive Union (cost=0.00..4.28 rows=101 width=4)

-> Result (cost=0.00..0.01 rows=1 width=4)

-> WorkTable Scan on s s_1 (cost=0.00..0.23 rows=10 width=4)

(6 rows)

Recursive CTEs

```
WITH RECURSIVE s AS (  
    SELECT 1 AS seq  
    UNION ALL  
    SELECT 2  
    UNION ALL  
    SELECT seq + 1  
    FROM s  
    WHERE seq <= 5  
) SELECT * FROM s;
```

Recursive CTEs

seq
1
2
2
3
3
4
4
5
5
6
6

(11 rows)

Recursive CTEs

```
WITH RECURSIVE s AS(  
    SELECT 1 AS seq, 'a' AS which  
    UNION ALL  
    SELECT 2, 'b'  
    UNION ALL  
    SELECT seq + 1, which  
    FROM s  
    WHERE seq <= 5  
) SELECT * FROM s;
```


Recursive CTEs

seq		which
1		a
2		b
2		a
3		b
3		a
4		b
4		a
5		b
5		a
6		b
6		a

(11 rows)

Recursive CTEs

seq	l	which
1	a	
2	b	
2	a, c	
3	b, c	
3	a, c, c	
4	b, c, c	
4	a, c, c, c	
5	b, c, c, c	
5	a, c, c, c, c	
6	b, c, c, c, c	
6	a, c, c, c, c, c	

(11 rows)

Recursive CTEs

```
WITH RECURSIVE s AS(  
    SELECT 1 AS seq, 'a' AS which  
    UNION ALL  
    SELECT 2, 'b'  
    UNION ALL  
    SELECT seq + 1, which || ', c'  
    FROM s  
    WHERE seq <= 5  
) SELECT * FROM s;
```

Homework!

```
CREATE TABLE recurse(  
  recurse_id serial PRIMARY KEY  
  , parent_id int REFERENCES recurse  
);
```

```
WITH i AS (INSERT INTO recurse(parent_id) VALUES(NULL)  
RETURNING *) INSERT INTO recurse(parent_id) SELECT  
recurse_id FROM i;
```

```
WITH i AS (INSERT INTO recurse(parent_id) VALUES(NULL)  
RETURNING *) INSERT INTO recurse(parent_id) SELECT  
recurse_id FROM i, generate_series(1,3) RETURNING *;
```

```
WITH i AS (INSERT INTO recurse(parent_id) VALUES(NULL)  
RETURNING *) INSERT INTO recurse(parent_id) SELECT 5 FROM i,  
generate_series(1,3) RETURNING *;
```


Homework!

```
CREATE TABLE recurse(  
  recurse_id serial PRIMARY KEY  
  , parent_id int REFERENCES recurse  
);
```

```
WITH i AS (INSERT INTO recurse(parent_id) VALUES(NULL)  
RETURNING *) INSERT INTO recurse(parent_id) SELECT  
recurse_id FROM i;
```

```
INSERT 0 1
```

How many rows are actually in the table now?

Homework!

```
WITH i AS (INSERT INTO recurse(parent_id) VALUES(NULL)
RETURNING *) INSERT INTO recurse(parent_id) SELECT
recurse_id FROM i, generate_series(1,3) RETURNING *;
```

recurse_id	parent_id
4	3
5	3
6	3

(3 rows)

Why is 4 the
starting recurse_id?

Questions?

Jim.Nasby@OpenSCG.com

<https://github.com/decibel/presentations/blob/master/CTEs.pdf>

Homework!

```
WITH i AS (INSERT INTO recurse(parent_id) VALUES(NULL)
RETURNING *) INSERT INTO recurse(parent_id) SELECT 5 FROM i,
generate_series(1,3) RETURNING *;
```

recurse_id	parent_id
8	5
9	5
10	5

(3 rows)

Homework!

```
WITH RECURSIVE l AS (  
  SELECT *, 1 AS level, recurse_id AS top_parent  
  FROM ...
```


Homework!

```
WITH RECURSIVE l AS (  
  SELECT *, 1 AS level, recurse_id AS top_parent  
    FROM recurse WHERE parent_id IS NULL  
  UNION ALL
```

Homework!

```
WITH RECURSIVE l AS (  
  SELECT *, 1 AS level, recurse_id AS top_parent  
    FROM recurse WHERE parent_id IS NULL  
  UNION ALL  
  SELECT ...  
    FROM l  
    JOIN recurse r ON ...
```

Homework!

```
WITH RECURSIVE l AS (  
  SELECT *, 1 AS level, recurse_id AS top_parent  
    FROM recurse WHERE parent_id IS NULL  
  UNION ALL  
  SELECT ...  
    FROM l  
    JOIN recurse r ON r.parent_id = l.recurse_id
```

Homework!

```
WITH RECURSIVE l AS (  
  SELECT *, 1 AS level, recurse_id AS top_parent  
    FROM recurse WHERE parent_id IS NULL  
  UNION ALL  
  SELECT r.recurse_id, r.parent_id, level+1, top_parent  
    FROM l  
    JOIN recurse r ON r.parent_id = l.recurse_id
```

Homework!

```
WITH RECURSIVE l AS (  
  SELECT *, 1 AS level, recurse_id AS top_parent  
    FROM recurse WHERE parent_id IS NULL  
  UNION ALL  
  SELECT r.recurse_id, r.parent_id, level+1, top_parent  
    FROM l  
    JOIN recurse r ON r.parent_id = l.recurse_id
```

recurse_id	parent_id	level	top_parent
1		1	1
3		1	3
7		1	7
2	1	2	1
4	3	2	3
5	3	2	3
6	3	2	3
8	5	3	3
9	5	3	3
10	5	3	3

(10 rows)